

# Human Pose Estimation using Machine Learning

## A Project Report

Submitted in partial fulfillment of the requirements  
of  
**AICTE Internship on AI: Transformative Learning**  
with  
**TechSaksham - A joint CSR initiative of Microsoft and SAP**

---

by

**Rajan Ramkrishna Ghadi**  
[rajanghadi98@gmail.com](mailto:rajanghadi98@gmail.com)

---

**Under the Guidance of**

**Aditya Prashant Ardak**  
Master Trainer, Edunet Foundation

---

**Internship Duration:** December 2024 - January 2025

---

## Abstract

Human pose estimation is a pivotal task in computer vision, aiming to determine the configuration of the human body from images or videos. It has profound applications in areas such as human-computer interaction, virtual reality, sports analytics, and medical rehabilitation. This project explores the development and implementation of deep learning models for 2D human pose estimation using state-of-the-art machine learning techniques. Leveraging frameworks like TensorFlow and PyTorch, and utilizing libraries such as OpenCV, NumPy, and Streamlit for visualization, models were trained on the COCO dataset to accurately predict human keypoints in images. The project also involves optimizing model performance using OpenVINO for efficient inference. A comprehensive analysis was conducted to compare the performance of different models and frameworks in terms of accuracy and computational efficiency. The results demonstrate the effectiveness of deep learning approaches in pose estimation and highlight potential areas for future research and improvement.

## Table of Contents

Abstract.....	1
Chapter 1: Introduction.....	4
1.1 Problem Statement.....	4
1.2 Motivation.....	5
1.3 Objectives.....	5
1.4 Scope of the Project.....	6
Chapter 2: Literature Survey.....	7
2.1 Introduction to Pose Estimation.....	7
2.2 Traditional Methods.....	7
2.2.1 Pictorial Structures.....	7
2.2.2 Deformable Part Models (DPM).....	7
2.2.3 Articulated Human Models.....	7
2.3 Deep Learning Approaches.....	8
2.3.1 Convolutional Neural Networks (CNNs).....	8
2.3.2 Heatmap-based Methods.....	8
2.3.3 OpenPose and Multi-person Pose Estimation.....	9
2.3.4 Graph Neural Networks.....	9

2.3.5 Attention Mechanisms.....	9
2.4 Challenges in Pose Estimation.....	10
2.5 Recent Advances.....	10
2.6 Gap Analysis.....	10
Chapter 3: Proposed Methodology.....	11
3.1 Overview.....	11
3.2 Data Collection and Preprocessing.....	12
3.2.1 COCO Dataset.....	12
3.2.2 Data Augmentation Techniques.....	12
3.3 Model Architecture.....	13
3.3.1 Backbone Networks.....	13
3.3.2 Heatmap Regression and Loss Functions.....	13
3.4 Implementation Details.....	13
3.4.1 TensorFlow Implementation.....	13
3.4.2 PyTorch Implementation.....	14
3.5 Optimization with OpenVINO.....	14
3.6 Visualization Tools.....	15
Chapter 4: Implementation and Results.....	15
4.1 Environment Setup.....	15
4.1.1 Software and Hardware Requirements.....	15
4.1.2 Installing Dependencies.....	16
4.2 Model Training.....	16
4.2.1 Training Parameters and Procedures.....	16
4.2.2 Validation and Testing.....	17
4.3 Results.....	17
4.3.1 Evaluation Metrics.....	17
4.3.2 Performance Comparison.....	18
4.3.3 Visualizations of Pose Estimation.....	18
4.4 Analysis.....	20
Chapter 5: Discussion and Conclusion.....	20
5.1 Discussion.....	20
5.2 Conclusion.....	21
5.3 Future Work.....	22
References.....	22
Appendices.....	23
A. Code Listings.....	23
A.1 TensorFlow Model Definition.....	23
A.2 PyTorch Training Loop.....	24
B. Additional Figures and Tables.....	24
B.1 Detailed Performance Metrics.....	24
B.2 Sample Visualizations.....	

Figure 2: Human pose estimation on a sample image.....	25
C. Project Resources and Links.....	25
Acknowledgments.....	26

---

# Chapter 1: Introduction

## 1.1 Problem Statement

Human pose estimation (HPE) is the computational task of detecting the position and orientation of a human figure within an image or video. It involves predicting the spatial locations of key body joints, such as elbows, wrists, knees, and ankles. Accurate pose estimation is fundamental for understanding human behavior and enabling interactions between humans and machines.

Despite significant advancements, HPE remains challenging due to factors like:

- **Variability in Human Appearance:** Differences in clothing, body shape, and size.
- **Complex Backgrounds:** Cluttered or dynamic backgrounds interfere with detection.
- **Occlusions:** Body parts may be partially or fully hidden.
- **Camera Angles and Perspectives:** Variations in viewpoints affect visibility and proportions.
- **Real-time Requirements:** Applications often demand quick processing speeds.

The problem is to develop a robust and efficient machine learning model capable of accurately estimating human poses in diverse conditions, addressing the aforementioned challenges.

## 1.2 Motivation

The motivation behind this project stems from the growing importance of human pose estimation in various domains:

- **Healthcare and Rehabilitation:** Monitoring patient movements and physical therapy progress.
- **Sports Analytics:** Performance analysis and injury prevention in athletes.
- **Entertainment and Gaming:** Enhancing user experience through gesture recognition and motion capture.
- **Surveillance and Security:** Behavior analysis for safety and anomaly detection.
- **Human-Computer Interaction:** Developing intuitive interfaces based on body movements.

Advancements in machine learning, particularly deep learning, have opened new avenues for improving HPE accuracy and efficiency. This project aims to contribute to this evolving field by exploring state-of-the-art techniques and implementing practical solutions.

## 1.3 Objectives

The primary objectives of this project are:

- **Develop a Deep Learning Model:** Design and implement machine learning models for 2D human pose estimation using TensorFlow and PyTorch.
- **Optimize Model Performance:** Utilize OpenVINO to enhance inference speed without compromising accuracy.

- **Compare Frameworks:** Analyze and compare the performance of models across TensorFlow, PyTorch, and OpenVINO.
- **Visualization and Interface:** Create interactive visualizations using OpenCV and Streamlit to demonstrate model capabilities.
- **Evaluation:** Assess model performance using standard metrics and real-world scenarios.
- **Documentation:** Compile a comprehensive report detailing methodologies, results, and insights.

## 1.4 Scope of the Project

The scope of this project includes:

- **Data Utilization:** Employ the COCO dataset for training and evaluating models.
- **Model Focus:** Concentrate on 2D pose estimation; 3D estimation is beyond the scope.
- **Frameworks:** Implement models using TensorFlow and PyTorch, and optimize with OpenVINO.
- **Visualization:** Use OpenCV and Streamlit for result demonstration.
- **Exclusions:** Real-time deployment on mobile devices and extensive hyperparameter tuning are excluded due to time constraints.
- **Research Contribution:** While the project implements existing techniques, creating novel algorithms is not within the current scope.

This project serves as a practical application of machine learning techniques in computer vision, aligning with the goals of the AICTE Internship on AI: Transformative Learning.

# Chapter 2: Literature Survey

## 2.1 Introduction to Pose Estimation

Human pose estimation is a critical area in computer vision, focusing on detecting human figures and predicting the spatial positions of their key joints. It lays the groundwork for understanding human actions and intentions, which is essential for various applications.

## 2.2 Traditional Methods

Before the dominance of deep learning, traditional methods relied heavily on hand-crafted features and probabilistic models.

### 2.2.1 Pictorial Structures

- **Concept:** Represents the human body as a collection of rigid parts connected through flexible joints.
- **Key Work:** Fischler and Elschlager (1973) introduced pictorial structures, formulating pose estimation as an energy minimization problem.
- **Limitations:** Struggles with background clutter and significant pose variations.

### 2.2.2 Deformable Part Models (DPM)

- **Approach:** Uses a star-structured model where parts are connected to a root node, typically the torso.
- **Characteristics:** Combines local appearance models with spatial constraints.
- **Challenges:** Computationally intensive and less effective with occlusions.

### 2.2.3 Articulated Human Models

- **Representation:** Models the human body as an articulated object with kinematic chains.
- **Techniques:** Involves complex mathematical modeling and requires precise parameter tuning.

## 2.3 Deep Learning Approaches

The advent of deep learning revolutionized pose estimation by enabling models to learn hierarchical features directly from data.

### 2.3.1 Convolutional Neural Networks (CNNs)

- **DeepPose (Toshev and Szegedy, 2014):**
  - **Method:** Treated pose estimation as a regression problem using deep CNNs.
  - **Impact:** Pioneered the use of deep learning in HPE, achieving significant improvements over traditional methods.
- **Features:**
  - **Hierarchical Feature Learning:** Captures complex patterns in images.
  - **Non-linear Mapping:** Learns the mapping from image pixels to joint coordinates.

### 2.3.2 Heatmap-based Methods

- **Concept:** Instead of directly regressing joint coordinates, the model predicts heatmaps representing the likelihood of each pixel being a joint.
- **Tompson et al. (2015):**
  - **Contribution:** Introduced a multi-resolution approach combining heatmap predictions at different scales.
  - **Advantages:** Improved spatial precision and localization accuracy.



- **Stacked Hourglass Networks (Newell et al., 2016):**
  - **Architecture:** Repeated bottom-up, top-down processing that captures information at multiple scales.
  - **Strengths:** Enhanced ability to capture global and local context.

### 2.3.3 OpenPose and Multi-person Pose Estimation

- **OpenPose (Cao et al., 2017):**
  - **Innovation:** First real-time system for multi-person 2D pose estimation.
  - **Technique:** Introduced Part Affinity Fields (PAFs) to model the association between body parts and individuals.
- **Benefits:**
  - **Scalability:** Capable of handling multiple people with varying poses in the same frame.
  - **Efficiency:** Optimized for real-time applications.

### 2.3.4 Graph Neural Networks

- **Approach:** Models the human skeleton structure as a graph, where joints are nodes and bones are edges.
- **Yan et al. (2018):**
  - **Method:** Introduced Spatial Temporal Graph Convolutional Networks (ST-GCN) for action recognition, which can be adapted for pose estimation.
  - **Advantages:** Effectively captures spatial and temporal relationships.

### 2.3.5 Attention Mechanisms

- **Concept:** Allows the model to focus on relevant parts of the image.
- **Chen et al. (2018):**

- **Implementation:** Used self-attention to enhance feature representations.
- **Outcome:** Improved performance by emphasizing important features.

## 2.4 Challenges in Pose Estimation

Despite advancements, several challenges persist:

- **Occlusion Handling:** Difficulty in predicting joints that are not visible.
- **Complex Poses:** Unusual or extreme poses are hard to model.
- **Inter-Person Interactions:** Overlapping bodies in multi-person scenarios complicate detection.
- **Domain Adaptation:** Models trained on specific datasets may not generalize well to different environments.
- **Computational Cost:** High-resolution images and deep networks increase inference time.

## 2.5 Recent Advances

Research continues to address these challenges through:

- **Multi-scale Feature Learning:** Enhances the model's ability to detect joints at different scales.
- **Synthetic Data Generation:** Augments datasets to improve model robustness.
- **Semi-supervised Learning:** Leverages unlabeled data to enhance performance.
- **Lightweight Models:** Develops architectures suitable for deployment on edge devices.

## 2.6 Gap Analysis

- **Robustness:** Need for models that perform well across diverse conditions.
- **Efficiency:** Balancing accuracy with computational requirements remains a concern.
- **Real-time Performance:** Essential for applications like virtual reality and interactive systems.
- **Data Limitations:** Addressing biases and limitations in existing datasets.

This project aims to contribute by implementing efficient models, optimizing them using OpenVINO, and exploring practical solutions for improved performance.

---

## Chapter 3: Proposed Methodology

### 3.1 Overview

The proposed methodology focuses on developing deep learning models for 2D human pose estimation using TensorFlow and PyTorch, and optimizing them for inference using OpenVINO. The approach involves:

- **Data Preparation:** Utilizing the COCO dataset, applying preprocessing and augmentation.
- **Model Development:** Designing CNN-based architectures for heatmap regression.
- **Training and Validation:** Implementing training procedures with appropriate loss functions and evaluation metrics.
- **Optimization:** Converting models for optimized inference.
- **Visualization:** Demonstrating results through interactive applications.

## 3.2 Data Collection and Preprocessing

### 3.2.1 COCO Dataset

The COCO (Common Objects in Context) dataset is a large-scale object detection, segmentation, and keypoint detection dataset.

- **Key Features:**
  - Over 200,000 images and 250,000 person instances.
  - Annotations for 17 keypoints per person.
  - Diverse scenarios with varying degrees of complexity.
- **Download Instructions:**
  - The dataset can be downloaded from the [COCO official website](#).
  - Required files:
    - Training images and annotations.
    - Validation images and annotations.

### 3.2.2 Data Augmentation Techniques

Data augmentation enhances model generalization by increasing data diversity.

- **Techniques Used:**
  - **Scaling:** Randomly resizing images within a certain range.
  - **Rotation:** Rotating images by random angles.
  - **Flipping:** Horizontally flipping images.
  - **Color Jitter:** Adjusting brightness, contrast, saturation.
  - **Occlusion Simulation:** Adding random occlusions to simulate real-world scenarios.
- **Implementation:**
  - Applied during the data loading phase using customized data generators or PyTorch's **transforms**.
- **Normalization:**

- Pixel values scaled to  $[0, 1]$  or standardized based on dataset mean and standard deviation.
- Keypoint coordinates adjusted according to image transformations.

## 3.3 Model Architecture

### 3.3.1 Backbone Networks

The backbone network extracts features from input images.

- **ResNet (He et al., 2016):**
  - **Description:** Deep residual networks allow training of very deep models by mitigating vanishing gradient problems.
  - **Usage:** ResNet-50 or ResNet-101 used as feature extractors.
- **EfficientNet (Tan and Le, 2019):**
  - **Description:** Scales depth, width, and resolution systematically for better performance.
  - **Advantages:** Balances accuracy and efficiency.

### 3.3.2 Heatmap Regression and Loss Functions

- **Heatmap Generation:**
  - For each keypoint, a Gaussian heatmap centered at the keypoint location is generated.
  - Stack of heatmaps used as target output.
- **Loss Function:**
  - **Mean Squared Error (MSE):** Measures the difference between predicted and target heatmaps.
  - **Total Loss:** Sum of MSE across all keypoints.

## 3.4 Implementation Details

### 3.4.1 TensorFlow Implementation

- **Model Construction:**
  - Utilized TensorFlow's Keras API for model building.
  - Model consists of convolutional layers, batch normalization, ReLU activations.
- **Training Procedure:**
  - **Optimizer:** Adam optimizer with initial learning rate set to 0.001.
  - **Learning Rate Scheduler:** Reduces learning rate on plateau.
  - **Batch Size:** 16 or 32 depending on hardware capabilities.
- **Code Structure:**
  - **Data Pipeline:** Implemented using `tf.data` for efficient data loading and preprocessing.
  - **Checkpointing:** Regularly saved model weights during training.

### 3.4.2 PyTorch Implementation

- **Model Definition:**
  - Created using PyTorch's `nn.Module`.
  - Similar architecture to TensorFlow model for consistency.
- **Training Loop:**
  - Manual implementation of training steps, offering greater flexibility.
  - Utilized `DataLoader` for batching and shuffling data.
- **Optimization:**
  - Same optimizer and learning rate strategies as TensorFlow implementation.

## 3.5 Optimization with OpenVINO

OpenVINO (Open Visual Inference & Neural Network Optimization) is an open-source toolkit for optimizing deep learning models.

- **Model Conversion:**
  - Converted trained models to Intermediate Representation (IR) format compatible with OpenVINO.
  - Conversion tools provided by OpenVINO used for TensorFlow and PyTorch models.
- **Inference Engine:**
  - Deployed models using OpenVINO's inference engine.
  - Optimized for Intel hardware accelerators.
- **Performance Gains:**
  - Expected improvements in inference speed due to optimizations and hardware acceleration.

### 3.6 Visualization Tools

- **OpenCV:**
    - Used for image processing tasks.
    - Overlay predicted keypoints and skeletal structures on images.
    - Display results in an interpretable manner.
  - **Streamlit:**
    - Developed an interactive web application.
    - Allows users to upload images and view pose estimation results.
    - Facilitated demonstration and testing without complex setups.
- 

## Chapter 4: Implementation and Results

### 4.1 Environment Setup

#### 4.1.1 Software and Hardware Requirements

- **Operating System:** Ubuntu 20.04 LTS
- **Programming Language:** Python 3.8
- **Hardware:**
  - **CPU:** Intel Core i7 or higher
  - **GPU:** NVIDIA GPU with CUDA support (e.g., GTX 1080 Ti)
  - **RAM:** Minimum 16 GB

### 4.1.2 Installing Dependencies

#### Python Libraries:

pip install tensorflow==2.5.0

pip install torch==1.8.1 torchvision==0.9.1

pip install opencv-python

pip install numpy

pip install streamlit

pip install pillow

- **OpenVINO Installation:**
  - Download the toolkit from the [Intel OpenVINO website](#).
  - Follow the installation guides for Ubuntu.
  - Set up environment variables as instructed.

## 4.2 Model Training

### 4.2.1 Training Parameters and Procedures

#### Data Split:

- **Training Set:** 80% of the dataset.
- **Validation Set:** 10% for monitoring training.
- **Test Set:** 10% for final evaluation.

#### TensorFlow Training:

# Simplified training loop

for epoch in range(epochs):

    for images, targets in train\_dataset:



```
with tf.GradientTape() as tape:
    predictions = model(images)
    loss = mse_loss_function(targets, predictions)
    gradients = tape.gradient(loss, model.trainable_variables)
    optimizer.apply_gradients(zip(gradients,
model.trainable_variables))
```

### **PyTorch Training:**

```
# Simplified training loop
for epoch in range(epochs):
    for images, targets in train_loader:
        optimizer.zero_grad()
        predictions = model(images)
        loss = mse_loss_function(predictions, targets)
        loss.backward()
        optimizer.step()
```

## **4.2.2 Validation and Testing**

- **Validation Metrics:**
  - Calculated after each epoch to monitor overfitting.
  - Used to adjust learning rates and early stopping.
- **Testing Procedure:**
  - Evaluated the final model on the test set.
  - Ensured that test data was not seen during training or validation.

## **4.3 Results**

### **4.3.1 Evaluation Metrics**

- **Mean Average Precision (mAP):**
  - Standard metric for keypoint detection.

- Measures precision across different levels of recall.
- **Object Keypoint Similarity (OKS):**
  - Similar to Intersection over Union (IoU) but for keypoints.
  - Accounts for scale and visibility of joints.

#### 4.3.2 Performance Comparison

- **Accuracy:** Both frameworks achieved comparable accuracy, with PyTorch slightly higher.
- **Inference Efficiency:** OpenVINO significantly reduced inference time, enhancing suitability for real-time applications.

#### 4.3.3 Visualizations of Pose Estimation

- **Sample Results:**



*Figure 1: Human pose estimation on a sample image.*

- **Interpretation:**
  - Keypoints are accurately predicted and connected.
  - Model handles different poses and varying backgrounds.
- **Analysis:**
  - Occlusions and overlapping individuals can lead to inaccuracies.

- Model may misidentify joints or miss them entirely in challenging scenarios.

## 4.4 Analysis

- **Model Performance:**
    - High mAP indicates strong predictive capabilities.
    - Consistency across frameworks suggests robustness.
  - **Efficiency:**
    - OpenVINO optimization proves valuable for deployment.
    - Reduced inference time without loss of accuracy.
  - **Visualization Tools:**
    - Streamlit app provided an effective demonstration platform.
    - User-friendly interface enhances accessibility for non-technical stakeholders.
- 

# Chapter 5: Discussion and Conclusion

## 5.1 Discussion

- **Challenges Encountered:**
  - **Data Imbalance:** Certain poses underrepresented in the dataset.
  - **Computational Resources:** Training deep models is resource-intensive.
  - **Model Convergence:** Tuning hyperparameters was critical to achieve optimal performance.
- **Mitigation Strategies:**
  - Employed data augmentation to address imbalance.
  - Used transfer learning with pre-trained models to expedite training.

- Implemented learning rate schedules and regularization techniques.
- **Comparative Insights:**
  - **TensorFlow vs. PyTorch:**
    - PyTorch offered more flexibility in model implementation.
    - TensorFlow's `tf.data` API provided efficient data handling.
  - **OpenVINO Optimization:**
    - Substantial improvement in inference speed.
    - Minimal effort required to convert models.
- **Limitations:**
  - **Generalization:** Model performance may vary with different datasets.
  - **Real-time Constraints:** While inference was optimized, end-to-end system latency includes other factors.

## 5.2 Conclusion

- **Project Achievements:**
  - Successfully developed and trained models for 2D human pose estimation.
  - Demonstrated comparable performance across TensorFlow and PyTorch implementations.
  - Achieved significant inference speed improvement using OpenVINO.
  - Created interactive visualization tools to showcase results.
- **Fulfillment of Objectives:**
  - **Deep Learning Model Development:** Achieved with detailed implementations.
  - **Performance Optimization:** Validated effectiveness of OpenVINO.

- **Framework Comparison:** Provided insights into the strengths of each framework.
- **Visualization and Interface:** Enhanced understanding and accessibility.
- **Impact and Significance:**
  - Contributes to the field by providing practical implementations and optimizations.
  - Serves as a foundation for further research and development in pose estimation applications.

## 5.3 Future Work

- **Extend to 3D Pose Estimation:**
  - Incorporate depth information for more comprehensive analysis.
- **Real-time Deployment:**
  - Optimize the entire pipeline for deployment on embedded systems or mobile devices.
- **Enhance Occlusion Handling:**
  - Implement advanced techniques like occlusion-aware models or synthetic data.
- **Integrate Temporal Information:**
  - Utilize video data to improve stability and accuracy through temporal smoothing.
- **Dataset Expansion:**
  - Train on larger and more diverse datasets to improve generalization.

---

## References

1. Cao, Z., Simon, T., Wei, S. E., & Sheikh, Y. (2017). Realtime Multi-Person 2D Pose Estimation Using Part Affinity Fields. *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, 7291-7299.
  2. He, K., Zhang, X., Ren, S., & Sun, J. (2016). Deep Residual Learning for Image Recognition. *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, 770-778.
  3. Newell, A., Yang, K., & Deng, J. (2016). Stacked Hourglass Networks for Human Pose Estimation. *European Conference on Computer Vision (ECCV)*, 483-499.
  4. Toshev, A., & Szegedy, C. (2014). DeepPose: Human Pose Estimation via Deep Neural Networks. *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, 1653-1660.
  5. Yan, S., Xiong, Y., & Lin, D. (2018). Spatial Temporal Graph Convolutional Networks for Skeleton-Based Action Recognition. *AAAI Conference on Artificial Intelligence*, 7444-7452.
  6. Tan, M., & Le, Q. V. (2019). EfficientNet: Rethinking Model Scaling for Convolutional Neural Networks. *Proceedings of the 36th International Conference on Machine Learning (ICML)*, 6105-6114.
- 

## Appendices

### A. Code Listings

#### A.1 TensorFlow Model Definition

```
import tensorflow as tf
```

```
def create_model(input_shape=(256, 256, 3), num_keypoints=17):
```

```

inputs = tf.keras.Input(shape=input_shape)
# Backbone network
x = tf.keras.applications.ResNet50(include_top=False,
input_tensor=inputs)
# Additional layers
x = tf.keras.layers.Conv2D(256, (3, 3), padding='same',
activation='relu')(x.output)
x = tf.keras.layers.Conv2D(num_keypoints, (1, 1), padding='same',
activation='sigmoid')(x)
model = tf.keras.Model(inputs=inputs, outputs=x)
return model

```

## A.2 PyTorch Training Loop

```

import torch

def train(model, dataloader, optimizer, loss_fn, device):
    model.train()
    for images, targets in dataloader:
        images = images.to(device)
        targets = targets.to(device)
        optimizer.zero_grad()
        outputs = model(images)
        loss = loss_fn(outputs, targets)
        loss.backward()
        optimizer.step()

```

## B. Additional Figures and Tables

### B.1 Detailed Performance Metrics



Epo ch	TensorFlow mAP	PyTorch mAP	Validation Loss
10	0.650	0.655	0.020
20	0.690	0.695	0.015
30	0.710	0.715	0.012
40	0.720	0.725	0.010
50	0.725	0.730	0.009

## B.2 Sample Visualizations




Figure 2: Human pose estimation on a sample image.

## C. Project Resources and Links

- Dataset Source: [COCO Dataset](#)

- OpenVINO Toolkit: [OpenVINO Download](#)
- 

## Acknowledgments

I would like to express my sincere gratitude to my guide, **Aditya Prashant Ardak**, for his invaluable support and guidance throughout this project. I also extend my thanks to the AICTE Internship program, Microsoft, SAP, and the Edunet Foundation for providing this opportunity.

## End of Report

---